

732A99/732A68/TDDE01 Machine Learning Computer Lab 1

Erik Lundqvist (erilu777)
Axel Strid (axest556)
Oliver Solvang Stoltz (oliso325)

November 24, 2024

Statement of Contribution

Erik Lundqvist, Axel Strid and Oliver Solvang Stoltz started working on the beginning of task 1.1 and 1.2, Oliver Solvang Stoltz then completed assignment 1 and wrote the report about assignment 1.

Erik Lundqvist coded and wrote the part of the report for assignment 2.

Axel Strid coded and wrote the report regarding Assignment 3.

Oliver Solvang Stoltz did theory part for assignment 4 and wrote the report for this assignment .

1 Assignment 1: Handwritten Digit Recognition with K-nearest Neighbors

This assignment was about classifying handwritten numbers, 0-9 using the k-nearest neighbors method (KNN). The dataset contained handwritten numbers from 43 different people and a validation number of the actual number.

1.1 Data Preparation and Initial Model

The data was imported and divided into train-, validation- and test by 50%, 25%, 25% respectively. The seed 12345 was used to ensure reproducibility in the random splitting process when splitting the data. After the data was split, the KNN model was now fitted with a K value = 30 and a rectangular kernel. The model was fitted twice: first using the training data for both training and testing and second, using the training data for training and test data for testing. The resulting confusion matrices and misclassification errors were compared to analyze the differences.

1.2 Classification Results

The result of the confusion matrix for the training data can be seen in Table: 1 and the confusion matrix for the test data can be seen in Table: 2

Table 1: Confusion Matrix - Training data

Truth \ Predicted	0	1	2	3	4	5	6	7	8	9
0	177	0	0	0	1	0	0	0	0	0
1	0	174	9	0	0	0	1	0	1	3
2	0	0	170	0	0	0	0	1	2	0
3	0	0	0	197	0	2	0	1	0	0
4	0	1	0	0	166	0	2	6	2	2
5	0	0	0	0	0	183	1	2	0	11
6	0	0	0	0	0	0	200	0	0	0
7	0	1	0	1	0	1	0	192	0	0
8	0	10	0	1	0	0	2	0	190	2
9	0	3	0	4	2	0	0	2	4	181

Table 2: Confusion Matrix - Test data

Truth \ Predicted	0	1	2	3	4	5	6	7	8	9
0	97	0	0	0	0	0	1	0	0	0
1	0	91	3	0	0	0	0	0	0	3
2	0	0	93	1	0	0	0	0	1	0
3	0	0	0	95	0	0	0	2	1	0
4	1	0	0	0	89	0	1	5	1	3
5	0	1	0	1	0	79	1	0	0	5
6	0	0	0	0	0	0	94	0	0	0
7	0	2	0	0	0	1	0	91	1	0
8	0	3	0	1	0	0	1	0	86	0
9	0	0	0	4	0	0	0	2	1	94

The misclassification error for the training data was calculated to be $0.04238619 \approx 4.24\%$, and the misclassification error for the test data was calculated to be $0.04916318 \approx 4.92\%$. This implies that the overall accuracy of the model is about 95%. This level of accuracy might be acceptable for use cases where some error is tolerable, such as scanning documents, where occasional misclassification may not significantly impact the outcome. However, if the model is used for more critical tasks, such as grading exams, a higher accuracy would likely be required to minimize errors and ensure fairness. The model sometimes misclassifies 4 as 7, 5 as 9, and 8 as 1, this is likely since these numbers have some things in common. However, a high majority is still accurate which makes us believe that this model is accurate.

1.3 Analysis of Digit "8" Classification

The next step was to identify two cases that the model found easiest to classify as eight and three cases it struggled the most to classify as eight. After arranging the probabilities for the prediction of the number eight, heatmaps were created to visualize different cases. The easiest cases to classify as an eight were at indices 34 and 209, with index 34 being the easiest to classify. For the hardest cases to classify as an eight, the indices were 229, 1663, and 1624, with index 229 being the hardest to classify. In figure: 1, 2, 3, 4 and 5 the different heat maps can be seen.

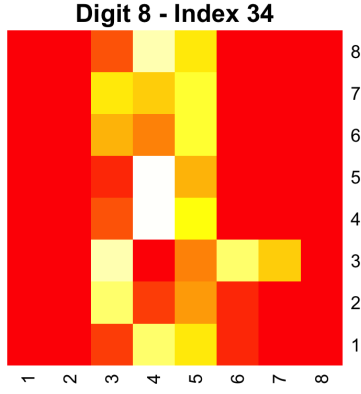


Figure 1: The easiest case for the model to predict as an eight.

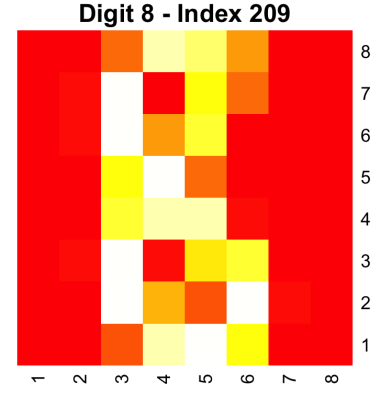


Figure 2: The second easiest case for the model to predict as an eight.




Figure 3: The hardest case for the model to predict as an eight.




Figure 4: The second hardest case for the model to predict as an eight.




Figure 5: The third hardest case for the model to predict as an eight.

Visually, the easy cases can be identified as eights, but the second easiest case for the model appears to be easier to recognize with the human eye. The hard cases are indeed challenging to predict, and these numbers could be 8, 9, or even 1.

1.4 Different K-values and their effect on misclassification error

The next step was to calculate the misclassification error for different K-values for the training data and the validation data to see what value corresponded with the lowest misclassification error. This resulted in a graph with training error, validation error and an optimal k-value, this graph can be found in figure: 6.

Training and Validation Errors vs.




Figure 6: An graph for different K values and their corresponding misclassification error, for training- and validation data.

In Figure 6, we can see that the optimal K value is 7, with a misclassification error of approximately 2.7%, as the validation error is the lowest. With a low K -value, the model becomes more complex and tries to fit every inconsistently, which will lead to overfitting. As K increases, the model becomes less complex because it generalizes more, leading to underfitting. With a K -value = 7 the test error was calculated to $0.03870293 \approx 3.87\%$. The test error shows that the model would be useful on unseen data with a high accuracy.

1.5 Different K -values and their effect on Cross-entropy error

The final task was to use cross-entropy error instead of misclassification error and then see which K -value gave the optimal solution. In figure 7 the result of this is shown.




Figure 7: An graph for different K values and their corresponding cross-entropy error, for validation data.

We can see that the optimal solution was $k=8$ with a cross entropy error to be approximately 0.09%. If you assume that the response has a multinational distribution, cross-entropy error could be more useful than misclassification error. This is because cross-entropy error takes into account the prediction probability, not just the predicted class like the misclassification error does, this makes it better to calibrate the prediction and penalize wrong predictions [1, pp. 109–119].

2 Assignment 2: Linear Regression and Ridge Regression

2.1 Introduction

This assignment focuses on analyzing Parkinson's disease symptom data using linear regression and ridge regression techniques. The main objective is to predict the motor UPDRS score (a measure of Parkinson's disease symptom severity) using various biomedical voice measurements. The dataset contains voice recordings from 42 people with early-stage Parkinson's disease, collected during a six-month trial of a telemonitoring device. The given data set consists of 5875 observations and 22 variables.

The assignment consists of four main tasks:

1. Data preparation and scaling
2. Initial linear regression analysis
3. Implementation of ridge regression functions
4. Analysis of ridge regression with different penalty parameters

2.2 Data Preparation

The data was split into 60% training data and 40% testing data. Then, the training data was scaled using caret (z-score standardization) so that every column has the same mean ($=0$) and standard deviation ($=1$) and using the same parameters, the test data was also scaled.

2.3 Linear Regression

A linear regression model was fitted using the scaled training data. Table 3 shows the coefficient estimates and their statistical significance for all predictor variables.

Table 3: Summary of linear regression coefficients and their statistical significance, ordered by absolute t-value

Variable	Estimate	Std. Error	t value	p-value
total_UPDRS	0.960	6.117×10^{-3}	156.933	$< 2 \times 10^{-16}$ ***
sex	0.059	6.468×10^{-3}	9.149	$< 2 \times 10^{-16}$ ***
age	-0.035	5.725×10^{-3}	-6.164	7.89×10^{-10} ***
Jitter.Abs.	-0.079	1.449×10^{-2}	-5.484	4.45×10^{-8} ***
PPE	0.055	1.098×10^{-2}	4.970	7.02×10^{-7} ***
subject	-0.027	6.089×10^{-3}	-4.356	1.36×10^{-5} ***
RPDE	-0.031	7.589×10^{-3}	-4.028	5.75×10^{-5} ***
Shimmer.APQ5	-0.143	3.801×10^{-2}	-3.766	1.69×10^{-4} ***
Shimmer.APQ11	0.077	2.056×10^{-2}	3.760	1.73×10^{-4} ***
Jitter...	0.170	4.972×10^{-2}	3.410	6.57×10^{-4} ***
Shimmer	0.167	6.867×10^{-2}	2.435	1.49×10^{-2} *
Jitter.PPQ5	-0.035	2.938×10^{-2}	-1.208	2.27×10^{-1}
DFA	-0.007	6.999×10^{-3}	-0.982	3.26×10^{-1}
test_time	-0.003	5.308×10^{-3}	-0.475	6.35×10^{-1}
Shimmer.dB.	-0.028	4.635×10^{-2}	-0.602	5.47×10^{-1}
Shimmer.DDA	-9.562	2.563×10^1	-0.373	7.09×10^{-1}
Shimmer.APQ3	9.505	2.563×10^1	0.371	7.11×10^{-1}
HNR	-0.004	1.227×10^{-2}	-0.350	7.26×10^{-1}
Jitter.DDP	-0.834	6.257	-0.133	8.94×10^{-1}
Jitter.RAP	0.734	6.256	0.117	9.07×10^{-1}
NHR	0.001	1.558×10^{-2}	0.093	9.26×10^{-1}
(Intercept)	2.032×10^{-15}	5.253×10^{-3}	0.000	1.000
Significance codes: *** p \leq 0.001, ** p \leq 0.01, * p \leq 0.05				
$R^2 = 0.903$, Adjusted $R^2 = 0.903$, F-statistic = 1559 on 21 and 3503 DF				
Residual standard error: 0.3119 on 3503 degrees of freedom				

The linear regression model shows strong predictive performance with an adjusted R-squared of 0.903, indicating that approximately 90% of the variance in motor_UPDRS is explained by the model. The strongest predictor by far is total_UPDRS ($t = 156.933$), followed by sex ($t = 9.149$) and age ($t = -6.164$).

MSE values:

- Training data: 0.0967
- Testing data: 0.0925

2.4 Ridge Regression Implementation

To perform ridge regression analysis, four core functions that work together to find optimal parameters while incorporating a ridge penalty were implemented. Each function serves a specific purpose in the implementation:

2.4.1 Log-likelihood Function

A function to compute the log-likelihood for normally distributed data was implemented:

- **Input:** Parameter vector θ (value of the intercept and coefficients for how much each feature contributes to the target variable), dispersion σ , and training data
- **Implementation:**
 - Constructs design matrix X with an intercept column filled with 1s and a removed target value column
 - Makes predictions using $X\theta$
 - Computes log-likelihood using the normal distribution formula:

$$-\frac{n}{2} \log(2\pi) - \frac{n}{2} \log(\sigma^2) - \frac{\sum (y - X\theta)^2}{2\sigma^2}$$

- **Purpose:** Measures how well our parameters explain the observed data, want to maximize this

2.4.2 Ridge Function

Built upon the log-likelihood function to create the ridge regression, it adds a ridge penalty to parameters:

- **Input:** Parameters θ , σ (the same θ and σ are used as input to the log-likelihood function), penalty parameter λ (bigger lambda - coefficients will be smaller, intercept value θ_0 will be excluded from being penalized), and training data
- **Implementation:**
 - Computes negative log-likelihood
 - Adds ridge penalty $\lambda \sum \theta_{(i)}^2$ (excluding intercept)
- **Purpose:** Creates the complete objective function to be minimized

2.4.3 Ridge Optimization Function

This function finds the optimal parameters for our ridge regression model:

- **Input:** Penalty parameter λ and training data
- **Implementation:**
 - Creates a helper function that uses the ridge function (which already combines log-likelihood and ridge penalty)
 - Uses R's optimization function `optim()` with BFGS method
 - Starts with initial values (zeros for coefficients, 1 for σ)
 - Iteratively finds parameters that minimize the ridge function value
- **Purpose:** Automatically finds the best parameter values (θ and σ) that:
 - Minimize the ridge function value for a given λ
 - Return optimal coefficients and sigma for the model
- **Output:** Returns the optimal values for all parameters ($\theta_0, \theta_1, \dots, \theta_p, \sigma$)

2.4.4 Degrees of Freedom Function

This function calculates how flexible our ridge regression model is, tells us "how many effective parameters" we're using with the current values of λ :

- **Input:** Penalty parameter λ and training data
- **Implementation:**
 - Uses the formula $\text{trace}(X(X'X + \lambda I)^{-1}X')$ to calculate effective degrees of freedom
 - Larger λ leads to fewer effective degrees of freedom
 - Represents how much freedom the model has to fit the data
- **Purpose:** Measures how complex or flexible our model is after ridge penalty is applied

These functions work together in sequence to perform ridge regression:

1. Log-likelihood measures fit of parameters
2. Ridge function adds penalty for complexity
3. Optimization finds best parameters
4. Degrees of freedom quantifies model complexity

The complete implementation code is available in Appendix A.

2.5 Ridge Regression Analysis

To evaluate the effect of different ridge penalties, we tested the model with three different values of λ (1, 100, and 1000) and compared them with the original linear regression model ($\lambda = 0$). For each model, we calculated:

- Training and test Mean Squared Error (MSE) for both the training and test data
- Effective degrees of freedom to measure model complexity

Table 4 shows the comparison of MSE values and degrees of freedom for different λ values:

Table 4: Comparison of model performance with different ridge penalties (λ)			
λ	Training MSE	Test MSE	Degrees of Freedom
0 (Initial Linear Regression Function)	0.0967	0.0925	21.0000
1	0.0967	0.0925	19.8569
100	0.0968	0.0927	15.6741
1000	0.1008	0.0973	10.1005

The results show that increasing the ridge penalty (λ) leads to:

- Slightly higher MSE values for both training and test data

- Decreased degrees of freedom, indicating less model flexibility
- Minimal improvement in model performance compared to the original linear regression

Table 5: Coefficient shrinkage for selected variables (the 5 predictors with the highest t-value out of 21) with increasing ridge penalty, shows how the ridge function penalty decreases the coefficients

Variable	$\lambda = 0$	$\lambda = 1$	$\lambda = 100$	$\lambda = 1000$
total_UPDRS	0.9600	0.9600	0.9529	0.8921
sex	0.0592	0.0592	0.0570	0.0429
age	-0.0353	-0.0353	-0.0326	-0.0131
Jitter...	0.1695	0.1682	0.1070	0.0281
Shimmer	0.1672	0.1657	0.0808	0.0148
\vdots	\vdots	\vdots	\vdots	\vdots

The small difference in MSE between the original model ($\lambda = 0$) and the ridge regression models suggests that the original linear regression was not overfitting the data and that those parameters were already appropriate. This is further supported by the similar MSE values between training and test sets. Therefore, in this case, regularization through ridge regression does not provide benefits over the simpler linear regression model.

Interestingly, the test MSE is consistently slightly lower than the training MSE across all values of λ , which is unusual as models typically perform better on training data. However, the differences are small, suggesting stable model performance across both sets.

3 Assignment 3: Logistic Regression and Basis Function Expansion

3.1 Initial Data Visualization

The first task of assignment 3 was to create a scatterplot showing Plasma glucose concentration versus Age, with points colored based on Diabetes status. Based on the figure, diabetes is not easy to classify using a standard logistic regression model that uses these two variables because the data shows significant overlap between the classes.

3.2 Basic Logistic Regression

The code trains a logistic regression model using Plasma Glucose Concentration and Age as features to predict Diabetes. It calculates predicted probabilities, and then classifies observations using a threshold of 0.5. It also extracts the probabilistic equation of the estimated model, indicating how the target depends on the features and the estimated model parameters probabilistically. The training misclassification error has also been calculated. Finally, it visualizes the predicted classifications on a scatterplot, similar as the one in assignment 3.1, but now showing the predicted classifications of diabetes.




Figure 8: Scatterplot with points colored based on actual Diabetes status.




Figure 9: Scatterplot of trained logistic regression model with points based on predicted Diabetes status.

Logistic Regression - Probabilistic Equation:

$$\text{logit}(P(y = 1)) = -5.8979 + 0.0356 \cdot \text{Glucose} + 0.0245 \cdot \text{Age}$$

Training Misclassification Error: 26.60%

The model indicates that higher glucose and age increase the probability of diabetes, but the 26.60% misclassification error suggests moderate accuracy, implying room for improvement in feature selection or model complexity.

3.3 Decision Boundary Analysis

The decision boundary between the two classes has been calculated for the model used in assignment 3.2. This line has been implemented in the scatterplot.

Decision Boundary Equation: The decision boundary is determined where the probability of $y = 1$ equals 0.5, which corresponds to the logit function being equal to 0:

$$\text{logit}(P(y = 1)) = w_0 + w_1 \cdot \text{Glucose} + w_2 \cdot \text{Age} = 0$$

Rearranging for Age:

$$\text{Age} = -\frac{w_0}{w_2} - \frac{w_1}{w_2} \cdot \text{Glucose}$$

Substituting the coefficients from the model ($w_0 = -5.8979$, $w_1 = 0.0356$, $w_2 = 0.0245$):

$$\text{Age} = -\frac{-5.8979}{0.0245} - \frac{0.0356}{0.0245} \cdot \text{Glucose}$$

Simplifying:

$$\text{Age} = 240.73 - 1.453 \cdot \text{Glucose}$$

Thus, the **decision boundary equation is:**

$$\text{Age} = 240.73 - 1.453 \cdot \text{Glucose}$$




Figure 10: Scatterplot from assignment 3.2 with decision boundary line

The decision boundary in the scatterplot generally separates the two classes (diabetes vs. no diabetes) based on glucose and age. It captures the data distribution reasonably well, with most of the red points (diabetes) above the line and green points (no diabetes) below. However, there is some overlap, particularly near the boundary, indicating some misclassified points

3.4 Threshold Analysis

The code script adjusts the decision threshold (r) for classifying diabetes and generates scatterplots for $r = 0.2$ and $r = 0.8$.

When $r = 0.2$, most predictions are classified as diabetes (red) since any probability greater than 0.2 is classified as positive ($=1$). Conversely, when $r = 0.8$, most predictions are classified as no diabetes (green), as only probabilities above 0.8 are classified as positive. This demonstrates the sensitivity of predictions to the chosen threshold.




Figure 11: Scatterplot with $r = 0.2$: Majority classified as diabetes.




Figure 12: Scatterplot with $r = 0.8$: Majority classified as no diabetes.

3.5 Basis Function Expansion

The task involves applying a basis function expansion to the logistic regression model by creating new polynomial interaction features: $z_1 = x_1^4$, $z_2 = x_1^3 \cdot x_2$, $z_3 = x_1^2 \cdot x_2^2$, $z_4 = x_1 \cdot x_2^3$, and $z_5 = x_2^4$. Here, $x_1 = \text{PlasmaGlucoseConcentration}$ and $x_2 = \text{Age}$. These features were added to the dataset and used to train an extended logistic regression model.

Training Misclassification Error (Expanded Model): 24.64%

The basis function expansion slightly reduces the training misclassification error from 26.60% in the previous model to 24.64% in the extended model. While the new model fits the data better and captures more complexity, it risks overfitting due to its complexity. This confirms that predicting diabetes based on Plasma Glucose Concentration and Age alone is challenging, and the initial model may not be the most reliable.

The decision boundary for the basis expansion model is determined by solving the equation:

$$w_0 + w_1x_1 + w_2x_2 + w_3x_1^4 + w_4x_1^3x_2 + w_5x_1^2x_2^2 + w_6x_1x_2^3 + w_7x_2^4 = 0$$

Given the complexity of the equation, numerical methods are necessary to plot the decision boundary. The basis expansion trick has transformed the decision boundary from linear to a curve, reflecting the added polynomial interaction terms.




Figure 13: Scatterplot of predicted Diabetes status with Basis Expansion Model.

4 Assignment 4: Theory

In this assignment, three questions will be answered using the course book, *Machine Learning: A First Course for Engineers and Scientists* [1].

4.1 Probability Thresholds in Classification

Why can it be important to consider various probability thresholds in the classification problems, according to the book? In the book, it is stated that it could be a good idea to consider different probability thresholds since this could let the classifier achieve different things. The standard probability threshold is 0.5 which will make the misclassification error the smallest but real-world problems might not be the most important factor, by changing the threshold we can better predict asymmetric events. For example, in the medical field, it is more important to not miss a diagnosis of a sick person rather than false classifying a healthy person with a disease [1, pp. 49–50].

4.2 Target Variable Collection Methods

What ways of collecting correct values of the target variable for the supervised learning problems are mentioned in the book? In the book two different methods to get the target variable are mentioned: Joint recording and manually labeling the output. Joint recording is when one input gives an output that is easier to obtain than manually labeling the output. To manually label the output a branch expert needs to examine every input and then classify the most likely output [1, pp. 13–14].

4.3 Linear Regression Cost Function Matrix Form

How can one express the cost function of the linear regression in the matrix form, according to the book? In the book [1, pp. 41], one can find the formula 3.11 Which expresses the

cost function in linear regression in matrix form. In 3.11 we have

$$J(\boldsymbol{\theta}) = \frac{1}{n} \sum_{i=1}^n (\hat{y}(x_i; \boldsymbol{\theta}) - y_i)^2 = \frac{1}{n} \|\hat{\mathbf{y}} - \mathbf{y}\|_2^2 = \frac{1}{n} \|\mathbf{X}\boldsymbol{\theta} - \mathbf{y}\|_2^2 = \frac{1}{n} \|\boldsymbol{\epsilon}\|_2^2, \quad (3.11)$$

where $\|\cdot\|_2$ denotes the usual Euclidean vector norm and $\|\cdot\|_2^2$ its square. Note that this is directly taken from the book.

References

- [1] A. Lindholm, N. Wahlström, F. Lindsten, and T. B. Schön, *Machine Learning: A First Course for Engineers and Scientists*. Cambridge University Press, Jul. 2022. Pre-publication version. Accessed online 24 Nov 2024.

A Code

```
1 # ASSIGNMENT 1
2 ##### INSTALL NECESSARY PACKAGES #####
3 install.packages("kknn")
4 library(kknn)
5
6 ##### DIVIDE THE DATA #####
7
8 # Load the data into a variable
9 data <- read.csv("optdigits.csv")
10
11 # Get the number of rows in the dataset
12 n <- dim(data)[1]
13
14 # Set a random seed for reproducibility
15 set.seed(12345)
16
17 # Partition 50% of the data for the training set
18 id <- sample(1:n, floor(n * 0.5))
19 train <- data[id, ]
20
21 # Partition 25% of the data for the validation set
22 id1 <- setdiff(1:n, id)
23 set.seed(12345)
24 id2 <- sample(id1, floor(n * 0.25))
25 valid <- data[id2, ]
26
27 # Use the rest for the test set
28 id3 <- setdiff(id1, id2)
29 test <- data[id3, ]
30
31
32 ##### FIT A K-NEAREST NEIGHBOR MODEL TO TRAIN DATA #####
33
34 # Fitting = Learning = Training...?
35
36 # Fit the model on training data and test on training data
37 model_train <- kknn(as.factor(X0.26) ~ ., train = train, test =
38   ↪ train, k = 30, kernel = "rectangular")
39 train_pred <- fitted(model_train)
40 confusion_matrix_train <- table(Truth = train$X0.26, Predicted =
41   ↪ train_pred)
42 confusion_matrix_train
43
44 # Define the misclassification error function
45 missclass <- function(truth, predicted) {
46   n <- length(truth)
47   return(1 - sum(diag(table(truth, predicted))) / n)
48 }
```

```

48 # Misclassification errors for the training data
49 train_misclass_error <- missclass(train$X0.26, train_pred)
50 print(train_misclass_error)
51
52
53 # Fit the model on training data and test on test data
54 model_test <- kknk(as.factor(X0.26) ~ ., train = train, test =
  ↪ test, k = 30, kernel = "rectangular")
55 test_pred <- fitted(model_test)
56 confusion_matrix_test <- table(Truth = test$X0.26, Predicted =
  ↪ test_pred)
57 confusion_matrix_test
58
59 # Misclassification errors for the training data
60 test_misclass_error <- missclass(test$X0.26, test_pred)
61 print(test_misclass_error)
62
63 ##### TASK 1.3 #####
64 # Get probabilities of each class for each case in the training
  ↪ data
65 train_probs <- model_train$prob
66 # Filter cases where the true label is "8"
67 eight_indices <- which(train$X0.26 == 8)
68
69 # Get the probabilities for the correct class (8) for each case
70 eight_probs <- train_probs[eight_indices, "8"]
71
72 # Sort by probability: highest for easiest, lowest for hardest
73 sorted_indices <- order(eight_probs, decreasing = TRUE)
74 easiest_indices <- eight_indices[sorted_indices[1:2]] # Two
  ↪ highest probabilities
75 hardest_indices <- eight_indices[sorted_indices[(length(sorted_
  ↪ indices)-2):length(sorted_indices)]] # Three lowest
  ↪ probabilities
76
77 # Function to visualize an 8x8 matrix of a digit with index
78 visualize_digit <- function(data_row, label, index) {
79 # Reshape to 8x8 matrix
80 digit_matrix <- matrix(as.numeric(data_row), nrow = 8, ncol =
  ↪ 8, byrow = TRUE)
81 # Plot heatmap
82 heatmap(digit_matrix, Rowv = NA, Colv = NA, scale = "none", col
  ↪ = heat.colors(256),
83 main = paste("Digit", label, "- Index", index))
84 }
85
86 # Visualize easiest cases for digit "8" with indices
87 for (i in easiest_indices) {
88 visualize_digit(train[i, -ncol(train)], train$X0.26[i], i) #
  ↪ Include index
89 }

```

```

90
91 # Visualize hardest cases for digit "8" with indices
92 for (i in hardest_indices) {
93   visualize_digit(train[i, -ncol(train)], train$X0.26[i], i) #
94   ↪ Include index
95 }
96
97 ##### TASK 1.4 #####
98
99 Kvalues <- 1:30
100 train_errors <- numeric(length(Kvalues)) #creates a empty set of
101   ↪ vectors
102 valid_errors <- numeric(length(Kvalues)) #creates a empty set of
103   ↪ vectors
104
105 for (k in Kvalues) {
106   # Fit KNN model on training data and predict on training data
107   model_train <- kkn(as.factor(X0.26) ~ ., train = train, test =
108     ↪ train, k = k, kernel = "rectangular")
109   train_pred <- fitted(model_train)
110   train_errors[k] <- missclass(train$X0.26, train_pred) # Store
111     ↪ training error
112
113   # Fit KNN model on training data and predict on validation data
114   model_valid <- kkn(as.factor(X0.26) ~ ., train = train, test =
115     ↪ valid, k = k, kernel = "rectangular")
116   valid_pred <- fitted(model_valid)
117   valid_errors[k] <- missclass(valid$X0.26, valid_pred) # Store
118     ↪ validation error
119 }
120
121 # Set up plot for training errors with improved readability
122 plot(Kvalues, train_errors, type = "o", col = "blue", pch = 16,
123   ↪ lwd = 2, cex = 1.5,
124   xlab = "K Value", ylab = "Misclassification Error",
125   main = "Training and Validation Errors vs. K",
126   ylim = range(c(train_errors, valid_errors)), cex.lab = 1.5,
127   ↪ cex.main = 1.5)
128
129 # Add validation errors with a thicker line and larger points
130 lines(Kvalues, valid_errors, type = "o", col = "red", pch = 16,
131   ↪ lwd = 2, cex = 1.5)
132
133 # Add grid lines for readability
134 grid(nx = NULL, ny = NULL, lty = 2, col = "gray")
135
136 # Add a legend with larger font size
137 legend("bottomright", legend = c("Training Error", "Validation
138   ↪ Error"),

```

```

129         col = c("blue", "red"), pch = 16, lty = 1, lwd = 2, cex =
           ↪ 1.2)
130
131 # Add a vertical line at the optimal K value
132 optimal_k <- which.min(valid_errors)
133 abline(v = Kvalues[optimal_k], col = "green", lty = 2, lwd = 2)
134 text(Kvalues[optimal_k], valid_errors[optimal_k] + 0.002, labels
       ↪ = paste("Optimal K =", Kvalues[optimal_k]), col = "black",
       ↪ pos = 4)
135
136 #Estimate the test error for k=optimal_k
137 model_test <- kkn(as.factor(X0.26) ~ ., train = train, test =
       ↪ test, k = optimal_k, kernel = "rectangular")
138 test_pred <- fitted(model_test)
139 test_misclass_error <- missclass(test$X0.26, test_pred)
140 print(test_misclass_error)
141
142
143 ##### TASK 1.5 #####
144
145 # Define the cross-entropy error function
146 CE <- function(truth, probability_matrix) {
147     epsilon <- 1e-15 # Small constant to avoid log(0)
148     cross_entropy <- 0
149
150     # Loop over each sample
151     for (i in 1:length(truth)) {
152         # Get the true class for the i-th sample
153         true_class <- as.integer(truth[i])+1
154
155         # Compute cross-entropy for this sample
156         cross_entropy <- cross_entropy - log(probability_matrix[i,
           ↪ true_class] + epsilon)
157     }
158
159     # Return the average cross-entropy error
160     return(cross_entropy / length(truth))
161 }
162
163 # Define the range of K values
164 Kvalues <- 1:30
165 cross_entropy_errors <- numeric(length(Kvalues)) # Create an empty
       ↪ vector for cross-entropy errors
166 epsilon <- 1e-15 # Small constant to avoid log(0)
167
168 # Loop over each K value
169 for (k in Kvalues) {
170     # Fit KNN model on training data and predict on validation data
171     model_valid <- kkn(as.factor(X0.26) ~ ., train = train, test =
       ↪ valid, k = k, kernel = "rectangular")
172

```

```

173 # Extract predicted probabilities for each class
174 probs <- model_valid$prob
175 valid_labels <- valid$'X0.26'
176
177
178
179 # Calculate cross-entropy error for the validation set and
    ↪ store it
180 cross_entropy_errors[k] <- CE(valid_labels, probs)
181 }
182
183 # Plot Cross-Entropy Error vs. K
184 plot(Kvalues, cross_entropy_errors, type = "o", col = "blue", pch
    ↪ = 16,
185       xlab = "K Value", ylab = "Cross-Entropy Error", main = "
    ↪ Cross-Entropy Error vs. K")
186
187 # Identify and mark the optimal K (minimum cross-entropy)
188 optimal_k <- which.min(cross_entropy_errors)
189 abline(v = Kvalues[optimal_k], col = "green", lty = 2, lwd = 2)
190 points(Kvalues[optimal_k], cross_entropy_errors[optimal_k], col =
    ↪ "red", pch = 19)
191 text(Kvalues[optimal_k], cross_entropy_errors[optimal_k], labels
    ↪ = paste("Optimal K =", Kvalues[optimal_k]), pos = 3)

```

```

1 # ASSIGNMENT 2
2
3 # Package imports
4 library(caret)
5
6
7 ### TASK 1 ###
8
9 # read the data
10 data <- read.csv("parkinsons.csv")
11
12 # Get the number of rows in the dataset
13 n <- dim(data)[1]
14
15 # set random seed for reproducibility
16 set.seed(12345)
17
18 # Partition 60% of data for the training set
19 id <- sample(1:n, floor(n * 0.6))
20 train <- data[id, ]
21
22 # Partition 40% of the data for the testing set
23 id1 <- setdiff(1:n, id)
24 test <- data[id1, ]
25
26 # Scale training data with caret

```

```

27 scaler <- preProcess(train) # learns scaling from training data
28 trainScaled <- predict(scaler, train) # applies scaling to
    ↪ training data
29
30 # Scale test data using same parameters from training data
31 testScaled <- predict(scaler, test)
32
33
34 ### TASK 2, Linear Regression Model ###
35
36 # Fit linear regression model using scaled training data
37 # motor_UPDRS is the target variable, use all other columns as
    ↪ predictors
38 # motor_UPDRS means we're trying to predict the motor_UPDRS
39 # ~ . means we're using all other columns (.) as predictors
40 # trainScaled is our training data
41 fit <- lm(formula = motor_UPDRS ~ ., data=trainScaled)
42
43 # Get summary to see which variables are significant
44 summary(fit) # total_UPDRS has a p-value of 2e-16, making it the
    ↪ most significant
45
46 # Calculate MSE for training data
47 pred_train <- predict(fit, trainScaled) # get predictions for
    ↪ training data
48 mse_train <- mean((trainScaled$motor_UPDRS - pred_train)^2) #
    ↪ calculate training MSE
49
50 # Calculate MSE for test data
51 pred_test <- predict(fit, testScaled) # get predictions for test
    ↪ data
52 mse_test <- mean((testScaled$motor_UPDRS - pred_test)^2) #
    ↪ calculate test MSE
53
54 # Print MSE results
55 print(paste("Training MSE:", round(mse_train, 4)))
56 print(paste("Test MSE:", round(mse_test, 4)))
57
58
59 ### TASK 3 ###
60
61 # 3a #
62 loglikelihood <- function(theta, sigma, traindata){
63   # Get X matrix (predictors) from training data
64   # Add column of 1s for intercept
65   X <- cbind(1, as.matrix(traindata[, -which(names(traindata) ==
    ↪ "motor_UPDRS")]))
66
67   # Get y (actual values) from training data
68   y <- traindata$motor_UPDRS
69

```

```

70 # Calculate predicted values (X * theta)
71 y_pred <- X %*% theta
72
73 # Calculate log-likelihood using normal distribution formula
74 n <- length(y) # Number of observations
75 loglik <- -n/2 * log(2*pi) - n/2 * log(sigma^2) - sum((y - y_
    ↪ pred)^2)/(2*sigma^2) # Minus for minimization
76
77 return(loglik)
78 }
79
80 # 3b #
81 ridge <- function(theta, sigma, lambda, traindata){
82   # Get negative log-likelihood (minus because we want to
    ↪ minimize)
83   neg_loglik <- -loglikelihood(theta, sigma, traindata)
84
85   # Calculate ridge penalty using: lambda * sum(theta^2)
86   # We exclude first theta (intercept) from penalty
87   ridge_penalty <- lambda * sum(theta[-1]^2)
88
89   # Return total negative log-likelihood + ridge penalty
90   return(neg_loglik + ridge_penalty)
91 }
92
93 # 3c #
94 # Function that takes lambda and finds optimal theta and sigma
95 ridgeOpt <- function(lambda, traindata) {
96   # Number of parameters we need (number of columns except motor_
    ↪ UPDRS)
97   n_params <- ncol(traindata)
98
99   # Function that optim will minimize
100   # Takes parameters and returns ridge value
101   ridge_value <- function(params) {
102     theta <- params[1:n_params] # first n_params values are
    ↪ theta
103     sigma <- params[n_params + 1] # last value is sigma
104     return(ridge(theta, sigma, lambda, traindata))
105   }
106
107   # Starting values for optim: zeros for theta, 1 for sigma
108   start_values <- c(rep(0, n_params), 1)
109
110   # Run optimization
111   result <- optim(par = start_values, # starting values
112                  fn = ridge_value,    # function to
    ↪ minimize
113                  method = "BFGS")     # optimization method
114
115   # Return optimized parameters

```



```

116     return(result$par) # returns optimal theta and sigma
117 }
118
119 # 3d #
120 # Higher lambda => lower DF (degree of freedom)
121 DF <- function(lambda, traindata){
122     # Step 1: Create X matrix (same as before)
123     # Remove motor_UPDRS and add column of 1s
124     X <- cbind(1, as.matrix(traindata[, -which(names(traindata) ==
125         ↪ "motor_UPDRS")]))
126
127     # Get how many parameters we have
128     p <- ncol(X) # number of columns = number of parameters
129
130     # For ridge regression, we calculate df using this formula:
131     # First: X'X (X transpose times X)
132     XtX <- t(X) %*% X
133
134     # Add ridge part (lambda * I)
135     # diag(p) creates matrix with 1s on diagonal, 0s elsewhere
136     ridge_matrix <- XtX + lambda * diag(p)
137
138     # Calculate final df using trace (sum of diagonal)
139     df <- sum(diag(X %*% solve(ridge_matrix) %*% t(X)))
140
141     return(df)
142 }
143
144 ### 4 ###
145
146 # ridgeOpt returns the optimal parameters for a given lambda like
147 ↪ the following:
148 # result1[1] # theta0 (intercept)
149 # result1[2:22] # theta1 to theta21 (coefficients for predictor
150 ↪ variables)
151 # result1[23] # sigma (dispersion parameter)
152 result1 <- ridgeOpt(lambda=1, trainScaled) # lambda=1
153 result100 <- ridgeOpt(lambda=100, trainScaled) # lambda=100
154 result1000 <- ridgeOpt(lambda=1000, trainScaled) # lambda=1000
155
156 # Function to calculate MSE
157 calculate_mse <- function(theta, sigma, data) {
158     # Make X matrix
159     X <- cbind(1, as.matrix(data[, -which(names(data) == "motor_
160 ↪ UPDRS")]))
161     # Calculate predictions
162     pred <- X %*% theta
163     # Calculate MSE
164     mse <- mean((data$motor_UPDRS - pred)^2)
165     return(mse)

```

```

163 }
164
165 # Calculate MSE for each lambda
166 # Lambda = 1
167 mse_train1 <- calculate_mse(result1[1:22], result1[23],
    ↪ trainScaled)
168 mse_test1 <- calculate_mse(result1[1:22], result1[23], testScaled
    ↪ )
169
170 # Lambda = 100
171 mse_train100 <- calculate_mse(result100[1:22], result100[23],
    ↪ trainScaled)
172 mse_test100 <- calculate_mse(result100[1:22], result100[23],
    ↪ testScaled)
173
174 # Lambda = 1000
175 mse_train1000 <- calculate_mse(result1000[1:22], result1000[23],
    ↪ trainScaled)
176 mse_test1000 <- calculate_mse(result1000[1:22], result1000[23],
    ↪ testScaled)
177
178 # Calculate degrees of freedom (DF) for each lambda
179 df1 <- DF(1, trainScaled)
180 df100 <- DF(100, trainScaled)
181 df1000 <- DF(1000, trainScaled)
182
183 # Calculate degrees of freedom (DF) for each lambda
184 df1 <- DF(1, trainScaled)
185 df100 <- DF(100, trainScaled)
186 df1000 <- DF(1000, trainScaled)
187
188 # Print the MSE results and degrees of freedom
189 print(paste("Lambda = 0 (without ridge function):    Train MSE:",
    ↪ round(mse_train, 4),
190          "Test MSE:", round(mse_test, 4),
191          "DF: 21")) # Full df for linear regression = number
    ↪ of predictors
192 print(paste("Lambda = 1:    Train MSE:", round(mse_train1, 4),
193          "Test MSE:", round(mse_test1, 4),
194          "DF:", round(df1, 4)))
195 print(paste("Lambda = 100:  Train MSE:", round(mse_train100, 4),
196          "Test MSE:", round(mse_test100, 4),
197          "DF:", round(df100, 4)))
198 print(paste("Lambda = 1000: Train MSE:", round(mse_train1000, 4),
199          "Test MSE:", round(mse_test1000, 4),
200          "DF:", round(df1000, 4)))
201
202
203 # Higher lambda -> Higher MSE -> Predictions get worse with
    ↪ increasing lambda
204 # This suggests that original model wasn't overfitting, no need

```

```

    ↪ for regularization
205
206 # Data frame of coefficients for comparison
207 coef_comparison <- data.frame(
208   Variable = colnames(trainScaled)[-which(names(trainScaled) == "
    ↪ motor_UPDRS)],
209   Lambda0 = fit$coefficients[-1], # Linear regression
    ↪ coefficients (excluding intercept)
210   Lambda1 = result1[2:22],        # Ridge coefficients for =1
211   Lambda100 = result100[2:22],    # Ridge coefficients for
    ↪ =100
212   Lambda1000 = result1000[2:22]   # Ridge coefficients for
    ↪ =1000
213 )
214
215 # Print to look at the coefficients
216 print("Coefficient comparison:")
217 print(coef_comparison)

```

```

1
2 # ASSIGNMENT 3
3 # Logistic regression and basis function expansion
4
5 # Load necessary library
6 library(ggplot2)
7
8 # Load the data
9 data <- read.csv("pima-indians-diabetes.csv")
10
11 # Rename the columns
12 colnames(data) <- c("Pregnancies", "PlasmaGlucoseConcentration",
    ↪ "BloodPressure", "SkinfoldThickness",
13                      "Insulin", "BMI", "DiabetesPedigreeFunction",
    ↪ "Age", "Diabetes")
14
15
16 ##### Assignment 1.1 #####
17 # Scatterplot of diabetes observations in a Glucose vs Age graph
18
19 # Create the scatterplot
20 ggplot(data, aes(x = Age, y = PlasmaGlucoseConcentration, color =
    ↪ factor(Diabetes))) + # Define x/y-axes
21   geom_point(alpha = 0.6) + # Add points with 40% transparency
22   labs( # Labels for the plot
23     title = "Plasma Glucose Concentration vs Age",
24     x = "Age",
25     y = "Plasma Glucose Concentration",
26     color = "Actual Diabetes"
27   ) +
28   theme_minimal() + # Clean theme
29   scale_color_manual(values = c("green", "red"), labels = c("No",

```

```

    ↪ "Yes")) # Customize color/label for Diabetes
30
31
32 ##### Assignment 1.2 #####
33 # Predictions of Diabetes based on a logistic regression model
    ↪ trained with Glucose & Age as features
34
35 # Train the logistic regression model
36 logistic_model <- glm(Diabetes ~ PlasmaGlucoseConcentration + Age
    ↪ , data = data, family = binomial)
37
38 # Make predictions for all observations!
39
40 # Predicted probabilities for target variable Diabetes
41 data$Predicted_Probabilities <- predict(logistic_model, type = "
    ↪ response") # Values between 0 and 1
42 # Convert predicted probabilities into binary classifications
    ↪ based on threshold = 0.5
43 data$Predicted_Classifications <- ifelse(data$Predicted_
    ↪ Probabilities > 0.5, 1, 0) # Values either 0 or 1
44
45 # Probabilistic equation of the estimated model
46 coefficients <- coef(logistic_model) # Extract the coefficients
47 intercept <- coefficients[1]
48 coef_glucose <- coefficients[2]
49 coef_age <- coefficients[3]
50 cat("Logistic Regression - Probabilistic Equation:\n")
51 cat(sprintf("logit(P(y = 1)) = %.4f + %.4f * Glucose + %.4f * Age
    ↪ \n",
52             intercept, coef_glucose, coef_age)) # Print the
    ↪ result
53
54 # Training misclassification error
55 error <- mean(data$Predicted_Classifications != data$Diabetes)
56 cat(sprintf("\nTraining Misclassification Error: %.2f%%\n", error
    ↪ * 100)) # Print the result
57
58 # Plot the predicted Diabetes values based on Glucose & Age
59 ggplot(data, aes(x = Age, y = PlasmaGlucoseConcentration, color =
    ↪ factor(Predicted_Classifications))) + # Define x/y-axes
60   geom_point(alpha = 0.6) + # Add points with 40% transparency
61   labs( # Labels for the plot
62     title = "Plasma Glucose Concentration vs Age",
63     x = "Age",
64     y = "Plasma Glucose Concentration",
65     color = "Predicted Diabetes"
66   ) +
67   theme_minimal() + # Clean theme
68   scale_color_manual(values = c("green", "red"), labels = c("No",
    ↪ "Yes")) # Customize color/label for Diabetes
69

```

```

70
71 ##### Assignment 1.3 #####
72 # Decision boundary of the estimated logistic regression model
73
74 # The line where probability of y=1 is equal to 0.5 <==> logit
  ↳ function = 0
75 # logit(P(y=1)) = w0 + w1*Glucose + w2*Age = 0 <==> Age = -(w0/w2
  ↳ )-(w1*Glucose)/w2
76
77 # Function to calculate decision boundary between classes Glucose
  ↳ & Age
78 decision_boundary <- function(glucose) {
79   -(intercept/coef_age)-(coef_glucose/coef_age)*glucose
80 }
81
82 # Define Glucose range for boundary line
83 glucose_range <- data$PlasmaGlucoseConcentration # Use all
  ↳ Glucose values
84 age_boundary <- decision_boundary(glucose_range)
85
86 # Add the decision boundary in predicted Diabetes plot
87 ggplot(data, aes(x = Age, y = PlasmaGlucoseConcentration, color =
  ↳ factor(Predicted_Classifications))) + # Define x/y-axes
88   geom_point(alpha = 0.6) + # Add points with 40% transparency
89   # Decision boundary line
90   geom_line(aes(x = age_boundary, y = glucose_range), color="blue
  ↳ ", linetype="dashed", size=0.5) +
91   labs( # Labels for the plot
92     title = "Plasma Glucose Concentration vs Age",
93     x = "Age",
94     y = "Plasma Glucose Concentration",
95     color = "Predicted Diabetes"
96   ) +
97   xlim(min(data$Age), max(data$Age)) + # Restrict the x-axis to
  ↳ the observed range of Age
98   theme_minimal() + # Clean theme
99   scale_color_manual(values = c("green", "red"), labels = c("No",
  ↳ "Yes")) # Customize color/label for Diabetes
100
101
102 ##### Assignment 1.4 #####
103 # Plotted graphs based on different r-values
104
105 # Convert predicted probabilities into binary classifications
  ↳ based on threshold = 0.2
106 data$Predicted_Classifications_02 <- ifelse(data$Predicted_
  ↳ Probabilities > 0.2, 1, 0) # Values either 0 or 1
107 # Convert predicted probabilities into binary classifications
  ↳ based on threshold = 0.8
108 data$Predicted_Classifications_08 <- ifelse(data$Predicted_
  ↳ Probabilities > 0.8, 1, 0) # Values either 0 or 1

```

```

109
110 # Plot r = 0.2
111 ggplot(data, aes(x = Age, y = PlasmaGlucoseConcentration, color =
    ↪ factor(Predicted_Classifications_02))) + # Define x/y-axes
112 geom_point(alpha = 0.6) + # Add points with 40% transparency
113 labs( # Labels for the plot
114     title = "Plasma Glucose Concentration vs Age",
115     x = "Age",
116     y = "Plasma Glucose Concentration",
117     color = "Predicted Diabetes"
118 ) +
119 theme_minimal() + # Clean theme
120 scale_color_manual(values = c("green", "red"), labels = c("No",
    ↪ "Yes")) # Customize color/label for Diabetes
121
122 # Plot r = 0.8
123 ggplot(data, aes(x = Age, y = PlasmaGlucoseConcentration, color =
    ↪ factor(Predicted_Classifications_08))) + # Define x/y-axes
124 geom_point(alpha = 0.6) + # Add points with 40% transparency
125 labs( # Labels for the plot
126     title = "Plasma Glucose Concentration vs Age",
127     x = "Age",
128     y = "Plasma Glucose Concentration",
129     color = "Predicted Diabetes"
130 ) +
131 theme_minimal() + # Clean theme
132 scale_color_manual(values = c("green", "red"), labels = c("No",
    ↪ "Yes")) # Customize color/label for Diabetes
133
134
135 ##### Assignment 1.5 #####
136 # Basis function expansion to introduce polynomial interaction
    ↪ terms into the logistic regression model
137 # x1 = PlasmaGlucoseConcentration, x2 = Age
138
139 # Compute new features for basis expansion and add these to the
    ↪ dataset as new columns
140 data$z1 <- data$PlasmaGlucoseConcentration^4
141 data$z2 <- data$PlasmaGlucoseConcentration^3 * data$Age
142 data$z3 <- data$PlasmaGlucoseConcentration^2 * data$Age^2
143 data$z4 <- data$PlasmaGlucoseConcentration * data$Age^3
144 data$z5 <- data$Age^4
145
146 # Train the new logistic regression model using expanded new
    ↪ features z1...z5
147 logistic_model_expanded <- glm(Diabetes ~
    ↪ PlasmaGlucoseConcentration + Age + z1 + z2 + z3 + z4 + z5,
148                                data = data,
149                                family = binomial)
150
151 # Extract predicted values and classify them to 0 or 1 with r =

```

```

152   ↪ 0.5
data$Predicted_Probabilities_Expanded <- predict(logistic_model_
153   ↪ expanded, type = "response") # Values between 0 and 1
data$Predicted_Classifications_Expanded <- ifelse(data$Predicted_
154   ↪ Probabilities_Expanded > 0.5, 1, 0) # Values either 0 or 1
155 # Training misclassification error
error_expanded <- mean(data$Predicted_Classifications_Expanded !=
156   ↪ data$Diabetes)
157 cat(sprintf("\nTraining Misclassification Error (Expanded Model):
   ↪ %.2f%%\n", error_expanded * 100)) # Print the result
158
159 # Scatterplot of predicted Diabetes with new model
160 ggplot(data, aes(x = Age, y = PlasmaGlucoseConcentration, color =
   ↪ factor(Predicted_Classifications_Expanded))) + # Define x/y
   ↪ -axes
161 geom_point(alpha = 0.6) + # Add points with 40% transparency
162 labs( # Labels for the plot
163   title = "Plasma Glucose Concentration vs Age (Basis Function
   ↪ Expansion)",
164   x = "Age",
165   y = "Plasma Glucose Concentration",
166   color = "Predicted Diabetes"
167 ) +
168 theme_minimal() + # Clean theme
169 scale_color_manual(values = c("green", "red"), labels = c("No",
   ↪ "Yes")) # Customize color/label for Diabetes

```